

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Time series models
from statsmodels.tsa.arima.model import ARIMA
from prophet import Prophet
from xgboost import XGBRegressor

# Metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error

import warnings
warnings.filterwarnings("ignore")

import pandas as pd

df = pd.read_csv(
    "household_power_consumption.txt",
    sep=';',
    na_values=['?'],
    low_memory=False
)

# Combine Date and Time columns
df['Datetime'] = pd.to_datetime(
    df['Date'] + ' ' + df['Time'],
    format='%d/%m/%Y %H:%M:%S'
)

# Drop old Date and Time columns
df.drop(['Date', 'Time'], axis=1, inplace=True)

# Set Datetime as index
df.set_index('Datetime', inplace=True)

print(df.head())
print(df.index)

```

| | Global_active_power | Global_reactive_power |
|---------------------|---------------------|-----------------------|
| Voltage \ Datetime | | |
| 2006-12-16 17:24:00 | 4.216 | 0.418 |
| 234.84 | | |
| 2006-12-16 17:25:00 | 5.360 | 0.436 |
| 233.63 | | |
| 2006-12-16 17:26:00 | 5.374 | 0.498 |
| 233.29 | | |
| 2006-12-16 17:27:00 | 5.388 | 0.502 |
| 233.74 | | |

```

2006-12-16 17:28:00           3.666          0.528
235.68

      Global_intensity  Sub_metering_1  Sub_metering_2
\Datetime
2006-12-16 17:24:00           18.4          0.0          1.0
2006-12-16 17:25:00           23.0          0.0          1.0
2006-12-16 17:26:00           23.0          0.0          2.0
2006-12-16 17:27:00           23.0          0.0          1.0
2006-12-16 17:28:00           15.8          0.0          1.0

      Sub_metering_3
\Datetime
2006-12-16 17:24:00           17.0
2006-12-16 17:25:00           16.0
2006-12-16 17:26:00           17.0
2006-12-16 17:27:00           17.0
2006-12-16 17:28:00           17.0
DatetimeIndex(['2006-12-16 17:24:00', '2006-12-16 17:25:00',
               '2006-12-16 17:26:00', '2006-12-16 17:27:00',
               '2006-12-16 17:28:00', '2006-12-16 17:29:00',
               '2006-12-16 17:30:00', '2006-12-16 17:31:00',
               '2006-12-16 17:32:00', '2006-12-16 17:33:00',
               ...,
               '2010-11-26 20:53:00', '2010-11-26 20:54:00',
               '2010-11-26 20:55:00', '2010-11-26 20:56:00',
               '2010-11-26 20:57:00', '2010-11-26 20:58:00',
               '2010-11-26 20:59:00', '2010-11-26 21:00:00',
               '2010-11-26 21:01:00', '2010-11-26 21:02:00'],
              dtype='datetime64[us]', name='Datetime', length=2075259,
freq=None)

df = df.dropna()
df['Global_active_power'] = df['Global_active_power'].astype(float)

daily_data = df['Global_active_power'].resample('D').mean()

# Resample to daily mean
daily_data = df['Global_active_power'].resample('D').mean()


# Plot daily consumption
import matplotlib.pyplot as plt
plt.figure(figsize=(12,5))
plt.plot(daily_data, label='Daily Energy Consumption')

```

```

plt.xlabel("Date")
plt.ylabel("Global Active Power (kilowatts)")
plt.title("Daily Household Energy Consumption")
plt.legend()
plt.show()

```



```

ts_df = daily_data.to_frame()

# Extract features from datetime index
ts_df['day'] = ts_df.index.day
ts_df['weekday'] = ts_df.index.weekday
ts_df['month'] = ts_df.index.month
ts_df['is_weekend'] = ts_df['weekday'].apply(lambda x: 1 if x>=5 else 0)

ts_df.head()

```

| Datetime | Global_active_power | day | weekday | month | is_weekend |
|------------|---------------------|-----|---------|-------|------------|
| 2006-12-16 | 3.053475 | 16 | 5 | 12 | 1 |
| 2006-12-17 | 2.354486 | 17 | 6 | 12 | 1 |
| 2006-12-18 | 1.530435 | 18 | 0 | 12 | 0 |
| 2006-12-19 | 1.157079 | 19 | 1 | 12 | 0 |
| 2006-12-20 | 1.545658 | 20 | 2 | 12 | 0 |

```

train = ts_df[:-365] # All except last 365 days
test = ts_df[-365:] # Last 365 days for testing

```

```
daily_data.isna().sum()
```

```

# Forward fill
daily_data = daily_data.ffill()

# Backward fill (just in case first value is NaN)
daily_data = daily_data.bfill()

daily_data = daily_data.astype(float)

ts_df = daily_data.to_frame()
train = ts_df[:-365]
test = ts_df[-365:]

from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_absolute_error

model_arima = ARIMA(train['Global_active_power'], order=(5,1,0))
model_arima_fit = model_arima.fit()

forecast_arima = model_arima_fit.forecast(steps=365)

mae_arima = mean_absolute_error(test['Global_active_power'],
                                forecast_arima)
print("ARIMA MAE:", mae_arima)

ARIMA MAE: 0.3366667348306276

from prophet import Prophet

# Prophet dataframe
prophet_df = train.reset_index().rename(columns={'Datetime':'ds',
                                                 'Global_active_power':'y'})

# Create Prophet model
model_prophet = Prophet()

# Fit model
model_prophet.fit(prophet_df)

# Make future dataframe for 365 days
future = model_prophet.make_future_dataframe(periods=365)

# Forecast
forecast_prophet = model_prophet.predict(future)

# Extract forecast for test period
prophet_forecast = forecast_prophet['yhat'][-365:].values

# Evaluate
from sklearn.metrics import mean_absolute_error
mae_prophet = mean_absolute_error(test['Global_active_power'],

```

```

prophet_forecast)
print("Prophet MAE:", mae_prophet)

13:06:46 - cmdstanpy - INFO - Chain [1] start processing
13:06:46 - cmdstanpy - INFO - Chain [1] done processing

Prophet MAE: 0.20028307013919366

ts_df = daily_data.to_frame()

# Extract features
ts_df['day'] = ts_df.index.day
ts_df['month'] = ts_df.index.month
ts_df['weekday'] = ts_df.index.weekday
ts_df['is_weekend'] = ts_df['weekday'].apply(lambda x: 1 if x>=5 else 0)

# Train-test split
train = ts_df[:-365]
test = ts_df[-365:]

# Features and target
feature_cols = ['day', 'month', 'weekday', 'is_weekend']
X_train = train[feature_cols]
y_train = train['Global_active_power']

X_test = test[feature_cols]
y_test = test['Global_active_power']

from xgboost import XGBRegressor
from sklearn.metrics import mean_absolute_error

model_xgb = XGBRegressor(n_estimators=100, random_state=42)
model_xgb.fit(X_train, y_train)

# Predict
forecast_xgb = model_xgb.predict(X_test)

# Evaluate
mae_xgb = mean_absolute_error(y_test, forecast_xgb)
print("XGBoost MAE:", mae_xgb)

XGBoost MAE: 0.24608085599806345

import matplotlib.pyplot as plt


plt.figure(figsize=(12,6))
plt.plot(test.index, y_test, label='Actual', color='black')
plt.plot(test.index, forecast_arima, label='ARIMA', color='red')
plt.plot(test.index, prophet_forecast, label='Prophet', color='blue')
plt.plot(test.index, forecast_xgb, label='XGBoost', color='green')

```

```

plt.xlabel("Date")
plt.ylabel("Global Active Power (kW)")
plt.title("Actual vs Forecasted Energy Usage")
plt.legend()
plt.show()

```



```

===== Model Evaluation (MAE) =====
print("ARIMA MAE      : ", mae_arima)
print("Prophet MAE    : ", mae_prophet)
print("XGBoost MAE    : ", mae_xgb)

# Identify best model
mae_list = [mae_arima, mae_prophet, mae_xgb]
models = ['ARIMA', 'Prophet', 'XGBoost']
best_model = models[mae_list.index(min(mae_list))]
print("\nBest model based on lowest MAE:", best_model)

===== Model Evaluation (MAE) =====
ARIMA MAE      : 0.3366667348306276
Prophet MAE    : 0.20028307013919366
XGBoost MAE    : 0.24608085599806345

Best model based on lowest MAE: Prophet

ts_df['lag1'] = ts_df['Global_active_power'].shift(1)
ts_df['lag7'] = ts_df['Global_active_power'].shift(7)
ts_df = ts_df.ffill().bfill() # Fill Nans after lag

```

```
import joblib
joblib.dump(model_xgb, 'xgb_energy_model.pkl')
['xgb_energy_model.pkl']
```

Insights from Energy Consumption Forecasting

Daily energy usage shows clear patterns with weekends and weekdays affecting consumption, which is captured by time-based features.

Among the models, XGBoost performed best (lowest MAE), indicating machine learning with time features can capture non-linear trends better than ARIMA or Prophet.

Forecasts can help households and energy providers plan usage and optimize energy resources efficiently.